

LAB Logbook

Lab 1

▼ Lab Logbook Requirement:

1) Create a vector using `np.arange`.

Determine the number of the vector elements using the following method: Take the last two digits from your SID. It should be from 00 to 99. If this number is 10 or more, it becomes the required number of the vector elements. If it is less than 10, add 100 to your number.

For example, if your SID is 2287467, and the last two digits are 67, which is greater than 10. The required number is 67. If your SID is 2287407, and the last two digits are 07, which is less than 10. The required number is 107.

Then,

2. Change matrix `a` to 2-d array with 1 row. Print the array. You should have the two sets of brackets for a 2-d array with one row.
3. Save it in another array. Print the array.
4. Check the shape attribute value.
5. Add the code and result to your Lab Logbook

```
[ ]:
[3]: sid=23368529%100
a = np.arange(sid) #sid = 2368529 as the last two digits of the sid is greater than 10
a1 = a.reshape(1, -1) #changing matrix a to 2-d array with 1 row
print(a1)
print('\n')
print(a1.shape)

[[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23
 24 25 26 27 28]]

(1, 29)
```

Lab 2

Lab Logbook Requirement:

1. Determine a number (n) equal to the last digit of your SID.
2. Group by "relationship" and "hours-per-week".
3. Reduce all "hours-per-week" column values in the original DataFrame by the value 'n'.
4. Group by "relationship" and reduced "hours-per-week".
5. Add the code and result to your Lab Logbook.

```
#Group by before reducing hours
```

```
Group_by_relationship = data.groupby(["relationship", "hours-per-week"])
Group_by_relationship.size()
```

```
relationship  hours-per-week
Husband      13              1
              40              4
              45              1
              80              1
Not-in-family 16              1
              40              2
              50              2
Own-child     30              1
Wife          40              2
dtype: int64
```

```
sid = 2368529%100
```

```
def func(x):
    return x - sid
```

```
data['hours-per-week'] = data['hours-per-week'].apply(func)
```

```
#Group by after reducing hours
```

```
Group_by_relationship = data.groupby(["relationship", "hours-per-week"])
Group_by_relationship.size()
```

```
relationship  hours-per-week
Husband      -16              1
              11              4
              16              1
              51              1
Not-in-family -13              1
              11              2
              21              2
Own-child     1              1
Wife          11              2
dtype: int64
```

Lab 3

Lab Logbook Requirement:

1) Draw a bicolour features interaction diagram between the columns with the numbers of the last and second to last digits of your SID, where:


```
# Column
---
1 Account length
2 Area code
3 International plan
4 Voice mail plan
5 Number vmail messages
6 Total day minutes
7 Total day calls
8 Total day charge
9 Total eve minutes
0 Total eve calls
```

In case these numbers are the same, then take the next number in order as another column number. For example, if your SID is 2287477, then you plot the bicolour diagram of the 7th and 8th columns. If your SID is 2287499, then the 9th and 0.

2. Add the code and result to your Lab Logbook.

```
[50]: #sid = 2368529 so plotted the bicolour diagram of the 2nd and 9th columns.
```

```
fig = plt.figure(figsize=(11,7))
plt.scatter(data['Area code'], data['Total eve minutes'], color = 'lr');
plt.xlabel('Area code');
plt.ylabel('Total eve minutes');
```



Lab 4

Lab Logbook Requirement:

1. Create your own Multi-layer Perceptron (MLP) with two hidden layers, where the first hidden layer cells' number equals the last three digits of your SID. The number of cells in the next hidden layer is approximately two times smaller. For example, if your SID is 2287167, the number of cells on the first hidden layer is 167, and on the second - 84. Take epochs=10. Leave other parameters the same as in the practical session.
2. Compile the model.
3. Train your MLP with the same datasets and demonstrate the received MAE.
4. Compare your MAE with the MAE of the MLP in the practical session.
5. Please only add to your Lab Logbook a print-screen of your MLP architecture using `model.summary()` and the resulting MAE.

¶

```
[47]: model = keras.Sequential([
      keras.layers.Dense(529, input_dim = 500, activation = tf.nn.relu, kernel_initializer = "normal"), #sid = 2368529
      keras.layers.Dense(265, activation = 'relu', kernel_initializer = "normal"),
      keras.layers.Dense(1)
    ])
print(model.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 529)	265,029
dense_17 (Dense)	(None, 265)	140,450
dense_18 (Dense)	(None, 1)	266

Total params: 405,745 (1.55 MB)
Trainable params: 405,745 (1.55 MB)
Non-trainable params: 0 (0.00 B)
None

```
[48]: model.compile(optimizer = "adam", loss = "mse", metrics = ["mae"])
```

```
[49]: history = model.fit(X_train, y_train, batch_size =10, epochs = 10, validation_split = 0.2, verbose = 1)
```

```
[47]: model = keras.Sequential([
    keras.layers.Dense(529, input_dim = 500, activation = tf.nn.relu, kernel_initializer = "normal"), #sid = 2368529
    keras.layers.Dense(265, activation = 'relu', kernel_initializer = "normal"),
    keras.layers.Dense(1)
])

print(model.summary())
```

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_16 (Dense)	(None, 529)	265,029
dense_17 (Dense)	(None, 265)	140,450
dense_18 (Dense)	(None, 1)	266

Total params: 405,745 (1.55 MB)
 Trainable params: 405,745 (1.55 MB)
 Non-trainable params: 0 (0.00 B)
 None

```
[48]: model.compile(optimizer = "adam", loss = "mse", metrics = ["mae"])
```

```
[49]: history = model.fit(X_train, y_train, batch_size = 10, epochs = 10, validation_split = 0.2, verbose = 1)
```

```
Epoch 1/10
2640/2640 — 20s 7ms/step - loss: 0.0356 - mae: 0.0442 - val_loss: 0.0107 - val_mae: 0.0936
Epoch 2/10
2640/2640 — 18s 7ms/step - loss: 2.1776e-04 - mae: 0.0115 - val_loss: 0.0076 - val_mae: 0.0794
Epoch 3/10
2640/2640 — 18s 7ms/step - loss: 1.9438e-04 - mae: 0.0108 - val_loss: 0.0048 - val_mae: 0.0614
Epoch 4/10
2640/2640 — 18s 7ms/step - loss: 1.1215e-04 - mae: 0.0081 - val_loss: 0.0013 - val_mae: 0.0300
Epoch 5/10
2640/2640 — 18s 7ms/step - loss: 9.0387e-05 - mae: 0.0071 - val_loss: 6.4911e-04 - val_mae: 0.0200
Epoch 6/10
2640/2640 — 22s 7ms/step - loss: 7.6753e-05 - mae: 0.0066 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 7/10
2640/2640 — 18s 7ms/step - loss: 6.4826e-05 - mae: 0.0061 - val_loss: 0.0012 - val_mae: 0.0310
Epoch 8/10
2640/2640 — 18s 7ms/step - loss: 5.9723e-05 - mae: 0.0058 - val_loss: 3.6299e-04 - val_mae: 0.0152
Epoch 9/10
2640/2640 — 18s 7ms/step - loss: 5.8251e-05 - mae: 0.0058 - val_loss: 4.3948e-04 - val_mae: 0.0172
Epoch 10/10
2640/2640 — 21s 7ms/step - loss: 5.0773e-05 - mae: 0.0055 - val_loss: 5.6606e-04 - val_mae: 0.0202
```

```
[58]: print("Mean absolute error: %.5f" % mae)
```

Mean absolute error: 0.01238

Lab 5

Lab Logbook Requirement:

1. Modify the practical session CNN model by reducing the convolutional core size to 5.

2. Change the batch_size to 50.

3. Also, change the size of the number of epochs, which is calculated by the formula:

$Z + Y$, if $Z = 0$

$10 + Y$, if $Z = 0$ and Y is not 0

10, if $Z = Y = 0$

, where your SID is: XXXXXZY

4. Leave other parameters the same as in the practical session.

5. Compile the model.

6. Train your CNN with the same datasets and demonstrate the received test MAE. Compare your MAE with the MAE of the CNN in the practical session.

7. Please only add a print-screen of your CNN architecture using `model.summary()` and the resulting MAE to your Lab Logbook.

¶

```
[39]: model = keras.Sequential([
      keras.layers.Conv1D(50,5, padding = 'same', input_shape= (50,5), activation=tf.nn.relu, kernel_initializer='normal'),
      keras.layers.MaxPooling1D(7),
      keras.layers.Conv1D(100,5,padding = 'same', activation = tf.nn.relu, kernel_initializer = "normal"),
      keras.layers.GlobalMaxPooling1D(),
      keras.layers.Dense(25, activation =tf.nn.relu, kernel_initializer = "normal"),
      keras.layers.Dense(2)
    ])
print(model.summary())
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv1d_4 (Conv1D)	(None, 50, 50)	1,300
max_pooling1d_2 (MaxPooling1D)	(None, 7, 50)	0
conv1d_5 (Conv1D)	(None, 7, 100)	25,100
global_max_pooling1d_2 (GlobalMaxPooling1D)	(None, 100)	0
dense_4 (Dense)	(None, 25)	2,525
dense_5 (Dense)	(None, 2)	52

Total params: 28,977 (113.19 KB)

Trainable params: 28,977 (113.19 KB)

Non-trainable params: 0 (0.00 B)

None

```
[40]: model.compile(optimizer = "adam", loss = "mse", metrics = ["mae"])
```

```
[41]: history = model.fit(X_train, y_train, batch_size=50, epochs=11, validation_split=0.2, verbose=1)
```

#sid = 2368529 where Z = 2 and Y = 9

```
Epoch 1/11
3520/3520 — 24s 6ms/step - loss: 0.0084 - mae: 0.0454 - val_loss: 9.3309e-04 - val_mae: 0.0202
Epoch 2/11
3520/3520 — 20s 6ms/step - loss: 7.5019e-04 - mae: 0.0186 - val_loss: 8.8981e-04 - val_mae: 0.0195
Epoch 3/11
3520/3520 — 20s 6ms/step - loss: 7.1502e-04 - mae: 0.0181 - val_loss: 8.3847e-04 - val_mae: 0.0186
Epoch 4/11
3520/3520 — 21s 6ms/step - loss: 6.9510e-04 - mae: 0.0177 - val_loss: 8.6578e-04 - val_mae: 0.0193
Epoch 5/11
3520/3520 — 21s 6ms/step - loss: 7.0214e-04 - mae: 0.0177 - val_loss: 8.2659e-04 - val_mae: 0.0186
Epoch 6/11
3520/3520 — 21s 6ms/step - loss: 6.7842e-04 - mae: 0.0175 - val_loss: 8.4197e-04 - val_mae: 0.0189
Epoch 7/11
3520/3520 — 21s 6ms/step - loss: 6.9582e-04 - mae: 0.0176 - val_loss: 8.3462e-04 - val_mae: 0.0187
Epoch 8/11
3520/3520 — 21s 6ms/step - loss: 7.0237e-04 - mae: 0.0176 - val_loss: 8.2371e-04 - val_mae: 0.0185
Epoch 9/11
3520/3520 — 43s 6ms/step - loss: 6.7800e-04 - mae: 0.0175 - val_loss: 8.7874e-04 - val_mae: 0.0196
Epoch 10/11
3520/3520 — 39s 6ms/step - loss: 6.9160e-04 - mae: 0.0175 - val_loss: 8.3649e-04 - val_mae: 0.0189
Epoch 11/11
3520/3520 — 21s 6ms/step - loss: 6.7548e-04 - mae: 0.0174 - val_loss: 8.6348e-04 - val_mae: 0.0194
```

```
[42]: mse,mae = model.evaluate(X_test, y_test, verbose=1)
print("Mean absolute error: %.5f" %mae)

936/936 — 3s 3ms/step - loss: 0.0012 - mae: 0.0237
Mean absolute error: 0.02526
```

Lab 6

Lab Logbook Requirement:

1. Plot the price chart of the part of the whole dataset 'High_Bid' and 'Low_Bid' prices using `iplot()` library.
2. The start point should equal the 5 last digits of your SID Number.
3. The time period (in minutes) should equal the 3 last digits of your SID Number.
4. Please only add a print-screen of your code and final graph to your Lab Logbook.



Lab 7

Lab Logbook Requirement:

1. Modify the practical session LSTM model parameter from 100 to be calculated using the formula:
 $ZY + 10$, where your SID is: XXXXXZY
2. Change the epochs to 10.
3. Change the patience to 3
4. Leave other parameters the same as in the practical session.
5. Compile the model.
6. Train your LSTM with the same datasets and demonstrate the received test MSE & MAE. Compare your test MSE & MAE with the MSE & MAE of the LSTM in the practical session.
7. Please only add to your Lab Logbook print-screens of:
 - your LSTM architecture using `model.summary()`,
 - the resulting test MSE & MAE and
 - MAE detailed graph

```
[50]: #sid =2368529 where Z=2 and Y=9
#ZY +10 = 29 + 10 = 39
```

```
model = keras.Sequential([
    keras.layers.LSTM(39, activation = 'relu', input_shape = (50, 18)),
    keras.layers.Dense(2)
])

print(model.summary())
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
lstm_3 (LSTM)	(None, 39)	9,048
dense_3 (Dense)	(None, 2)	80

Total params: 9,128 (35.66 KB)
 Trainable params: 9,128 (35.66 KB)
 Non-trainable params: 0 (0.00 B)
 None

```
[51]: model.compile(optimizer = "adam", loss = "mse", metrics = ["mae"])
```

```
[52]: es = EarlyStopping(monitor='val_loss', mode='min', patience=3, verbose=1)
mc = ModelCheckpoint('best_model_LSTM_GOLD.keras', monitor='val_loss', mode='min', verbose=1, save_best_only=True)
```

```
[53]: history = model.fit(X_train, y_train, batch_size = 20, epochs = 10, validation_split = 0.1, shuffle = True, verbose =1, callbacks = [es,mc])
```

```
Epoch 1/10
1212/1213 ----- 0s 16ms/step - loss: 0.1920 - mae: 0.1231
Epoch 1: val_loss improved from inf to 0.00005, saving model to best_model_LSTM_GOLD.keras
1213/1213 ----- 25s 18ms/step - loss: 0.1918 - mae: 0.1230 - val_loss: 4.9154e-05 - val_mae: 0.0049
Epoch 2/10
1212/1213 ----- 0s 17ms/step - loss: 4.8671e-05 - mae: 0.0052
Epoch 2: val_loss improved from 0.00005 to 0.00003, saving model to best_model_LSTM_GOLD.keras
1213/1213 ----- 21s 17ms/step - loss: 4.8660e-05 - mae: 0.0052 - val_loss: 2.9932e-05 - val_mae: 0.0039
Epoch 3/10
1211/1213 ----- 0s 17ms/step - loss: 3.3265e-05 - mae: 0.0045
Epoch 3: val_loss did not improve from 0.00003
1213/1213 ----- 21s 17ms/step - loss: 3.3264e-05 - mae: 0.0045 - val_loss: 6.5819e-05 - val_mae: 0.0068
Epoch 4/10
1212/1213 ----- 0s 19ms/step - loss: 3.3284e-05 - mae: 0.0045
Epoch 4: val_loss did not improve from 0.00003
1213/1213 ----- 44s 19ms/step - loss: 3.3282e-05 - mae: 0.0045 - val_loss: 6.8866e-05 - val_mae: 0.0069
Epoch 5/10
1213/1213 ----- 0s 18ms/step - loss: 3.2081e-05 - mae: 0.0045
Epoch 5: val_loss improved from 0.00003 to 0.00002, saving model to best_model_LSTM_GOLD.keras
1213/1213 ----- 22s 18ms/step - loss: 3.2080e-05 - mae: 0.0045 - val_loss: 2.0336e-05 - val_mae: 0.0033
Epoch 6/10
1212/1213 ----- 0s 16ms/step - loss: 3.7082e-05 - mae: 0.0048
Epoch 6: val_loss did not improve from 0.00002
1213/1213 ----- 21s 17ms/step - loss: 3.7076e-05 - mae: 0.0048 - val_loss: 3.4457e-05 - val_mae: 0.0047
Epoch 7/10
1213/1213 ----- 0s 15ms/step - loss: 3.2006e-05 - mae: 0.0045
Epoch 7: val_loss did not improve from 0.00002
1213/1213 ----- 19s 15ms/step - loss: 3.2006e-05 - mae: 0.0045 - val_loss: 2.4473e-05 - val_mae: 0.0044
Epoch 8/10
1211/1213 ----- 0s 17ms/step - loss: 2.4711e-05 - mae: 0.0039
Epoch 8: val_loss improved from 0.00002 to 0.00001, saving model to best_model_LSTM_GOLD.keras
1213/1213 ----- 21s 17ms/step - loss: 2.4712e-05 - mae: 0.0039 - val_loss: 1.3461e-05 - val_mae: 0.0030
Epoch 9/10
1211/1213 ----- 0s 20ms/step - loss: 2.2413e-05 - mae: 0.0037
Epoch 9: val_loss did not improve from 0.00001
1213/1213 ----- 45s 21ms/step - loss: 2.2414e-05 - mae: 0.0037 - val_loss: 4.0357e-05 - val_mae: 0.0050
Epoch 10/10
1213/1213 ----- 0s 19ms/step - loss: 2.0940e-05 - mae: 0.0036
Epoch 10: val_loss improved from 0.00001 to 0.00001, saving model to best_model_LSTM_GOLD.keras
1213/1213 ----- 24s 20ms/step - loss: 2.0939e-05 - mae: 0.0036 - val_loss: 9.3247e-06 - val_mae: 0.0026
```

```
[67]: scores = LSTM_saved_best_model.evaluate(X_test, y_test, verbose=1)
```

```
94/94 ----- 1s 6ms/step - loss: 9.4234e-06 - mae: 0.0026
```

```
[68]: scores
```

```
[68]: [8.694913049112074e-06, 0.0024753068573772907]
```

```
[69]: print("Mean squared error (mse): %.9f " % (scores[0]))
```

```
Mean squared error (mse): 0.000008695
```


```
[70]: print("Mean absolute error (mae): %.9f " % (scores[1]))
```

```
Mean absolute error (mae): 0.002475307
```

```
[71]: history_dict = history.history
```

```
mae_values = history_dict['mae']
val_mae_values = history_dict['val_mae']
```

```
epochs = range(1, len(mae_values) + 1)
plt.figure(num=1, figsize=(15,7))
plt.plot(epochs, mae_values, 'b', label='Training Mean Absolute Error(MAE)')
plt.plot(epochs, val_mae_values, marker='o', markeredgcolor='red', markerfacecolor='yellow', label='Validation Mean Absolute Error(MAE)')
plt.xlabel('Epochs', size=18)
plt.ylabel('Mean Absolute Error(MAE)', size=18)
plt.legend()
plt.show()
```

Lab 8

```
[1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
%matplotlib inline

from matplotlib import*
import seaborn as sns
```

```
[2]: bid=pd.read_csv("XAGUSD_5 Mins_Bid_2023.01.01_2023.06.30.csv")
ask=pd.read_csv("XAGUSD_5 Mins_Ask_2023.01.01_2023.06.30.csv")
```

```
[3]: mlp=bid.merge(ask, left_on='Time (UTC)', right_on='Time (UTC)', how='outer')
mlp
```

```
[3]:
```

	Time (UTC)	Open_x	High_x	Low_x	Close_x	Volume_x	Open_y	High_y	Low_y	Close_y	Volume_y
0	2023.01.02 23:00:00	24.036	24.059	24.017	24.059	0.3480	24.102	24.125	24.083	24.125	0.3510
1	2023.01.02 23:05:00	24.064	24.130	24.064	24.092	1.8458	24.094	24.188	24.094	24.141	1.1550
2	2023.01.02 23:10:00	24.094	24.098	23.972	23.977	0.9030	24.143	24.148	24.022	24.027	0.8820
3	2023.01.02 23:15:00	23.977	23.980	23.938	23.980	0.8940	24.026	24.028	23.986	24.028	0.8940
4	2023.01.02 23:20:00	23.978	24.024	23.976	24.023	0.8880	24.026	24.073	24.023	24.073	0.9044
...
35215	2023.06.30 20:35:00	22.752	22.752	22.741	22.746	0.2752	22.782	22.782	22.771	22.776	1.8002
35216	2023.06.30 20:40:00	22.736	22.751	22.736	22.746	0.1410	22.775	22.781	22.775	22.776	0.7568
35217	2023.06.30 20:45:00	22.751	22.751	22.746	22.746	0.0690	22.781	22.781	22.776	22.776	0.4212
35218	2023.06.30 20:50:00	22.746	22.761	22.736	22.756	0.2160	22.776	22.797	22.774	22.786	0.6836
35219	2023.06.30 20:55:00	22.756	22.766	22.701	22.745	0.3270	22.786	22.811	22.786	22.811	0.3712

35220 rows x 11 columns

```
[4]: mlp.columns = ['time', 'open_bid', 'high_bid', 'low_bid', 'close_bid', 'volume_bid', 'open_ask', 'high_ask', 'low_ask', 'close_ask', 'volume_ask']
```

```
[5]: mlp.head()
```

```
[5]:
```

	time	open_bid	high_bid	low_bid	close_bid	volume_bid	open_ask	high_ask	low_ask	close_ask	volume_ask
0	2023.01.02 23:00:00	24.036	24.059	24.017	24.059	0.3480	24.102	24.125	24.083	24.125	0.3510
1	2023.01.02 23:05:00	24.064	24.130	24.064	24.092	1.8458	24.094	24.188	24.094	24.141	1.1550
2	2023.01.02 23:10:00	24.094	24.098	23.972	23.977	0.9030	24.143	24.148	24.022	24.027	0.8820
3	2023.01.02 23:15:00	23.977	23.980	23.938	23.980	0.8940	24.026	24.028	23.986	24.028	0.8940
4	2023.01.02 23:20:00	23.978	24.024	23.976	24.023	0.8880	24.026	24.073	24.023	24.073	0.9044

```
[6]: file_obj2 = open('mlp.csv', 'w')
mlp.to_csv('mlp.csv', encoding='utf-8', index=False)
file_obj2.close()
```

```
[7]: new=pd.read_csv('mlp.csv', low_memory=False, sep=',')
```

```
[8]: new.describe()
org=new.drop(['open_ask', 'high_ask', 'low_ask', 'close_ask'],axis=1)
org.shape
org.head(3)
```

```
[8]:
```

	time	open_bid	high_bid	low_bid	close_bid	volume_bid	volume_ask
0	2023.01.02 23:00:00	24.036	24.059	24.017	24.059	0.3480	0.351
1	2023.01.02 23:05:00	24.064	24.130	24.064	24.092	1.8458	1.155
2	2023.01.02 23:10:00	24.094	24.098	23.972	23.977	0.9030	0.882

```
[9]: org['time'] = pd.to_datetime(org['time'])
org.head(3)
```


```
[9]:
```

	time	open_bid	high_bid	low_bid	close_bid	volume_bid	volume_ask
0	2023-01-02 23:00:00	24.036	24.059	24.017	24.059	0.3480	0.351
1	2023-01-02 23:05:00	24.064	24.130	24.064	24.092	1.8458	1.155
2	2023-01-02 23:10:00	24.094	24.098	23.972	23.977	0.9030	0.882

```
[10]: plt.figure(figsize=(12,6))
plt.plot(org['close_bid'])
plt.title('SILVER close price, 2023')
plt.xlabel('N minutes')
plt.ylabel('Close Price')
plt.show()
```



```
[11]: plt.figure(figsize=(12,6))
plt.plot(org['volume_bid'])
plt.title('SILVER close Volume, 2023')
plt.xlabel('N minutes')
plt.ylabel('Close Volume')
plt.show()
```



```
[12]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split # Added import
from sklearn.preprocessing import MinMaxScaler
from keras.models import Sequential
from keras.layers import Dense

# Assuming 'org' is your dataset

# Normalize the selected columns
scaler = MinMaxScaler()

columns_to_normalize = ['open_bid', 'high_bid', 'low_bid', 'close_bid', 'volume_bid', 'volume_ask']
org[columns_to_normalize] = scaler.fit_transform(org[columns_to_normalize])

# Shifting the 'high_bid' and 'low_bid' columns as an example
org['high_bid_shifted_next'] = org['high_bid'].shift(-1) # Shift by 1 for next value
org['low_bid_shifted_next'] = org['low_bid'].shift(-1)

# Drop the Last row as the shifted column will have NaN value
org = org.dropna()

# Separate the input (X) and output (y) matrices
X = org[['open_bid', 'high_bid', 'low_bid', 'volume_bid', 'volume_ask', 'high_bid_shifted_next', 'low_bid_shifted_next']].values
y = org['close_bid'].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Define the MLP model
model = Sequential()

# Add Dense Layers (MLP Layers)
model.add(Dense(units=64, activation='relu', input_dim=X_train.shape[1])) # Input Layer
model.add(Dense(units=32, activation='relu')) # Hidden Layer
model.add(Dense(units=1)) # Output Layer for regression (single value prediction)

# Compile the model with MSE Loss and MAE as metrics
model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])

# Train the model
history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_test, y_test))

# Get the final MSE and MAE values
final_mse = history.history['loss'][-1]
final_mae = history.history['mae'][-1]

print(f"Final MSE: {final_mse}")
print(f"Final MAE: {final_mae}")

# Plotting the Loss (MSE) and MAE over epochs
plt.figure(figsize=(12, 6))

# Plot MSE
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train MSE')
plt.plot(history.history['val_loss'], label='Validation MSE', linestyle='--')
plt.title('Model Loss (MSE)')
plt.xlabel('Epochs')
plt.ylabel('MSE')
plt.legend()

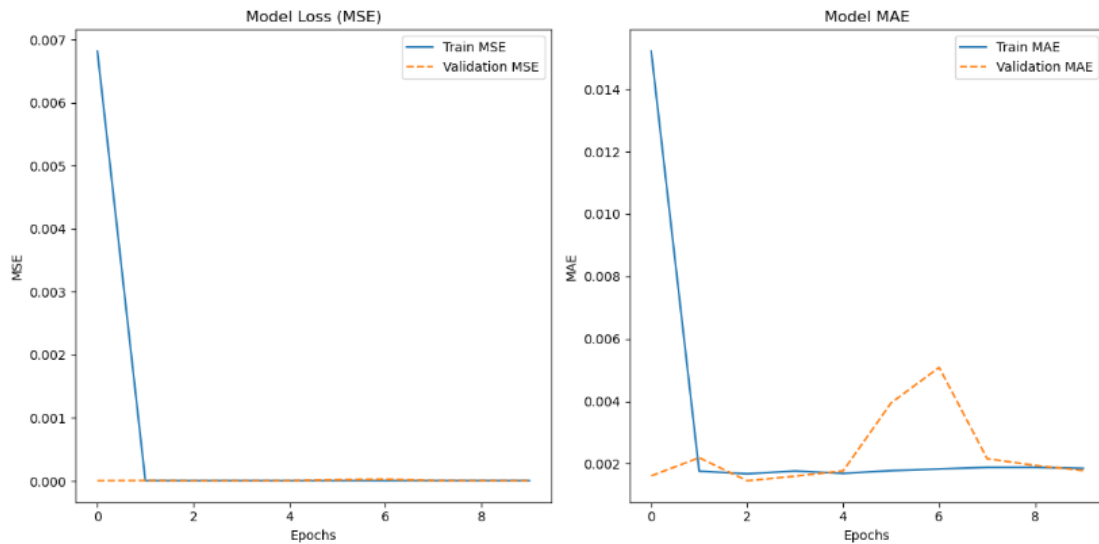
# Plot MAE
plt.subplot(1, 2, 2)
plt.plot(history.history['mae'], label='Train MAE')
plt.plot(history.history['val_mae'], label='Validation MAE', linestyle='--')
plt.title('Model MAE')
plt.xlabel('Epochs')
plt.ylabel('MAE')
plt.legend()

plt.tight_layout()
plt.show()
```

```

Epoch 1/10
881/881 10s 6ms/step - loss: 0.0361 - mae: 0.0631 - val_loss: 5.7601e-06 - val_mae: 0.0016
Epoch 2/10
881/881 4s 5ms/step - loss: 6.2113e-06 - mae: 0.0018 - val_loss: 7.9547e-06 - val_mae: 0.0022
Epoch 3/10
881/881 4s 5ms/step - loss: 5.7862e-06 - mae: 0.0017 - val_loss: 4.5911e-06 - val_mae: 0.0015
Epoch 4/10
881/881 4s 4ms/step - loss: 5.7670e-06 - mae: 0.0017 - val_loss: 5.0067e-06 - val_mae: 0.0016
Epoch 5/10
881/881 4s 4ms/step - loss: 5.2190e-06 - mae: 0.0016 - val_loss: 5.4796e-06 - val_mae: 0.0018
Epoch 6/10
881/881 4s 4ms/step - loss: 6.2169e-06 - mae: 0.0019 - val_loss: 1.9841e-05 - val_mae: 0.0040
Epoch 7/10
881/881 4s 4ms/step - loss: 6.2499e-06 - mae: 0.0019 - val_loss: 3.0096e-05 - val_mae: 0.0051
Epoch 8/10
881/881 4s 5ms/step - loss: 6.5881e-06 - mae: 0.0019 - val_loss: 7.3881e-06 - val_mae: 0.0022
Epoch 9/10
881/881 4s 4ms/step - loss: 7.2862e-06 - mae: 0.0021 - val_loss: 6.2551e-06 - val_mae: 0.0019
Epoch 10/10
881/881 4s 4ms/step - loss: 5.8475e-06 - mae: 0.0018 - val_loss: 5.4987e-06 - val_mae: 0.0018
Final MSE: 6.0957518144277856e-06
Final MAE: 0.0018504991894587874

```



Lab 10

Lab Logbook Requirement:

- Plot 4 graphs:
 - Precision during training graph
 - More detailed Precision graph
 - Training accuracy graph
 - More detailed Accuracy graph

Precision during training

```
[280]: import numpy as np
import pandas as pd
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score
import matplotlib.pyplot as plt

daily_data = pd.read_csv("XAUUSD_Daily_Ask_2024.01.01_2024.06.30.csv")
daily_data["Time (UTC)"] = pd.to_datetime(daily_data["Time (UTC)"])
daily_data.set_index("Time (UTC)", inplace=True)
daily_data = daily_data[["Close"]]

daily_data["Target"] = (daily_data["Close"].shift(-1) > daily_data["Close"]).astype(int)
daily_data.dropna(inplace=True)
daily_data["Close"] = (daily_data["Close"] - daily_data["Close"].mean()) / daily_data["Close"].std()

sequence_length = 10
X, y = [], []
for i in range(len(daily_data) - sequence_length):
    X.append(daily_data["Close"].iloc[i:i + sequence_length].values)
    y.append(daily_data["Target"].iloc[i + sequence_length])
X = np.array(X)
y = np.array(y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
```

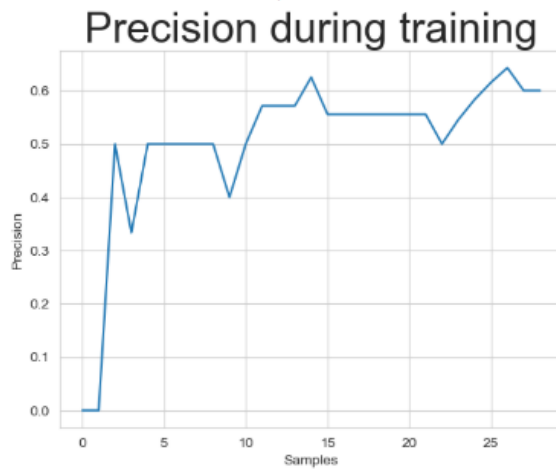
```
model = Sequential([
    LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'),
    Dense(1, activation='sigmoid')
])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
model.fit(X_train, y_train, epochs=20, batch_size=8, verbose=1, validation_split=0.2)
```

```
y_pred = (model.predict(X_test) > 0.5).astype(int)
precision_values = [precision_score(y_test[i+1], y_pred[i+1]) for i in range(len(y_test))]
```

```
plt.plot(precision_values)
plt.xlabel('Samples')
plt.ylabel('Precision')
plt.title('Precision during training')
plt.show()
```

```
Epoch 1/20
12/12 — 8s 100ms/step - accuracy: 0.4964 - loss: 0.6892 - val_accuracy: 0.5417 - val_loss: 0.6971
Epoch 2/20
12/12 — 0s 15ms/step - accuracy: 0.5676 - loss: 0.6906 - val_accuracy: 0.5417 - val_loss: 0.6903
Epoch 3/20
12/12 — 0s 15ms/step - accuracy: 0.6327 - loss: 0.6751 - val_accuracy: 0.5417 - val_loss: 0.6885
Epoch 4/20
12/12 — 0s 15ms/step - accuracy: 0.6105 - loss: 0.6727 - val_accuracy: 0.5000 - val_loss: 0.6823
Epoch 5/20
12/12 — 0s 14ms/step - accuracy: 0.6093 - loss: 0.6765 - val_accuracy: 0.5000 - val_loss: 0.6825
Epoch 6/20
12/12 — 0s 14ms/step - accuracy: 0.5913 - loss: 0.6666 - val_accuracy: 0.4583 - val_loss: 0.6805
Epoch 7/20
12/12 — 0s 15ms/step - accuracy: 0.6574 - loss: 0.6594 - val_accuracy: 0.4583 - val_loss: 0.6818
Epoch 8/20
12/12 — 0s 14ms/step - accuracy: 0.6020 - loss: 0.6804 - val_accuracy: 0.4583 - val_loss: 0.6837
Epoch 9/20
12/12 — 0s 14ms/step - accuracy: 0.5893 - loss: 0.6843 - val_accuracy: 0.4583 - val_loss: 0.6855
Epoch 10/20
12/12 — 0s 23ms/step - accuracy: 0.6747 - loss: 0.6499 - val_accuracy: 0.4583 - val_loss: 0.6873
Epoch 11/20
12/12 — 0s 16ms/step - accuracy: 0.5657 - loss: 0.6932 - val_accuracy: 0.4583 - val_loss: 0.6892
Epoch 12/20
12/12 — 0s 16ms/step - accuracy: 0.6586 - loss: 0.6717 - val_accuracy: 0.4583 - val_loss: 0.6929
Epoch 13/20
12/12 — 0s 14ms/step - accuracy: 0.6440 - loss: 0.6404 - val_accuracy: 0.4583 - val_loss: 0.6938
Epoch 14/20
12/12 — 0s 14ms/step - accuracy: 0.6297 - loss: 0.6202 - val_accuracy: 0.4583 - val_loss: 0.6952
Epoch 15/20
12/12 — 0s 14ms/step - accuracy: 0.5458 - loss: 0.6875 - val_accuracy: 0.4583 - val_loss: 0.6987
Epoch 16/20
12/12 — 0s 14ms/step - accuracy: 0.6407 - loss: 0.6522 - val_accuracy: 0.4583 - val_loss: 0.7001
Epoch 17/20
12/12 — 0s 14ms/step - accuracy: 0.5890 - loss: 0.6839 - val_accuracy: 0.5000 - val_loss: 0.7023
Epoch 18/20
12/12 — 0s 14ms/step - accuracy: 0.6303 - loss: 0.6463 - val_accuracy: 0.5417 - val_loss: 0.7073
Epoch 19/20
12/12 — 0s 14ms/step - accuracy: 0.6384 - loss: 0.6702 - val_accuracy: 0.4583 - val_loss: 0.7147
Epoch 20/20
12/12 — 0s 14ms/step - accuracy: 0.6079 - loss: 0.6543 - val_accuracy: 0.4583 - val_loss: 0.7219
1/1 — 1s 591ms/step
```



```
[282]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from tensorflow.keras.callbacks import Callback
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score

daily_data = pd.read_csv("XAUUSD_Daily_Ask_2024.01.01_2024.06.30.csv")
daily_data['Time (UTC)'] = pd.to_datetime(daily_data['Time (UTC)'])
daily_data.set_index('Time (UTC)', inplace=True)
daily_data = daily_data[['Close']]

daily_data['Target'] = (daily_data['Close'].shift(-1) > daily_data['Close']).astype(int)
daily_data.dropna(inplace=True)
daily_data['Close'] = (daily_data['Close'] - daily_data['Close'].mean()) / daily_data['Close'].std()

sequence_length = 10
X, y = [], []
for i in range(len(daily_data) - sequence_length):
    X.append(daily_data['Close'].iloc[i:i + sequence_length].values)
    y.append(daily_data['Target'].iloc[i + sequence_length])
X = np.array(X)
y = np.array(y)

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_val = X_val.reshape((X_val.shape[0], X_val.shape[1], 1))

class PrecisionCallback(Callback):
    def __init__(self):
        self.train_precision = []
        self.val_precision = []
```

```

def on_epoch_end(self, epoch, logs=None):
    y_train_pred = (self.model.predict(X_train) > 0.5).astype(int)
    train_precision = precision_score(y_train, y_train_pred)
    self.train_precision.append(train_precision)

    y_val_pred = (self.model.predict(X_val) > 0.5).astype(int)
    val_precision = precision_score(y_val, y_val_pred)
    self.val_precision.append(val_precision)

model = Sequential([
    LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'),
    Dense(1, activation='sigmoid')
])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

precision_callback = PrecisionCallback()
model.fit(X_train, y_train, epochs=20, batch_size=8, verbose=1, validation_data=(X_val, y_val), callbacks=[precision_callback])

epochs = range(1, len(precision_callback.train_precision) + 1)
plt.plot(epochs, precision_callback.train_precision, label='Training Precision', color='blue')
plt.plot(epochs, precision_callback.val_precision, label='Validation Precision', color='yellow', marker='o')
plt.xlabel('Epochs')
plt.ylabel('Precision')
plt.title('More detailed precision graph')
plt.legend()
plt.show()

```

```

Epoch 1/20
4/4 ————— 1s 198ms/step - accuracy: 0.6091 - loss: 0.6
1/1 ————— 0s 65ms/step
15/15 ————— 1s 204ms/step - accuracy: 0.5720 - loss: 0.6900 - val_accuracy: 0.5862 - val_loss: 0.6891
Epoch 2/20
4/4 ————— 0s 6ms/step p - accuracy: 0.5556 - loss: 0.685
1/1 ————— 0s 72ms/step
15/15 ————— 1s 46ms/step - accuracy: 0.5554 - loss: 0.6857 - val_accuracy: 0.5172 - val_loss: 0.6894
Epoch 3/20
4/4 ————— 0s 8ms/step p - accuracy: 0.5978 - loss: 0.681
1/1 ————— 0s 73ms/step
15/15 ————— 1s 52ms/step - accuracy: 0.5953 - loss: 0.6819 - val_accuracy: 0.5172 - val_loss: 0.6895
Epoch 4/20
4/4 ————— 0s 7ms/step p - accuracy: 0.5181 - loss: 0.688
1/1 ————— 0s 62ms/step
15/15 ————— 1s 45ms/step - accuracy: 0.5219 - loss: 0.6883 - val_accuracy: 0.5517 - val_loss: 0.6879
Epoch 5/20
4/4 ————— 0s 6ms/step p - accuracy: 0.5645 - loss: 0.667
1/1 ————— 0s 63ms/step
15/15 ————— 1s 45ms/step - accuracy: 0.5658 - loss: 0.6685 - val_accuracy: 0.5517 - val_loss: 0.6885
Epoch 6/20
4/4 ————— 0s 7ms/step p - accuracy: 0.5862 - loss: 0.671
1/1 ————— 0s 68ms/step
15/15 ————— 1s 47ms/step - accuracy: 0.5856 - loss: 0.6718 - val_accuracy: 0.5517 - val_loss: 0.6886
Epoch 7/20
4/4 ————— 0s 6ms/step p - accuracy: 0.5729 - loss: 0.681
1/1 ————— 0s 66ms/step
15/15 ————— 1s 44ms/step - accuracy: 0.5732 - loss: 0.6814 - val_accuracy: 0.5517 - val_loss: 0.6884
Epoch 8/20
4/4 ————— 0s 7ms/step ep - accuracy: 0.5728 - loss: 0.677
1/1 ————— 0s 66ms/step
15/15 ————— 1s 49ms/step - accuracy: 0.5747 - loss: 0.6776 - val_accuracy: 0.5517 - val_loss: 0.6903
Epoch 9/20
4/4 ————— 0s 5ms/step p - accuracy: 0.5767 - loss: 0.656
1/1 ————— 0s 65ms/step
15/15 ————— 1s 45ms/step - accuracy: 0.5767 - loss: 0.6576 - val_accuracy: 0.5517 - val_loss: 0.6909
Epoch 10/20
4/4 ————— 0s 5ms/step p - accuracy: 0.5584 - loss: 0.664
1/1 ————— 0s 65ms/step
15/15 ————— 1s 46ms/step - accuracy: 0.5596 - loss: 0.6654 - val_accuracy: 0.5517 - val_loss: 0.6895
Epoch 11/20
4/4 ————— 0s 6ms/step p - accuracy: 0.5833 - loss: 0.665
1/1 ————— 0s 65ms/step
15/15 ————— 1s 44ms/step - accuracy: 0.5835 - loss: 0.6662 - val_accuracy: 0.5517 - val_loss: 0.6883
Epoch 12/20
4/4 ————— 0s 5ms/step p - accuracy: 0.5861 - loss: 0.6585
1/1 ————— 0s 64ms/step

```



```

15/15 ----- 1s 44ms/step - accuracy: 0.5785 - loss: 0.6666 - val_accuracy: 0.5172 - val_loss: 0.6876
Epoch 13/20
4/4 ----- 0s 6ms/step p - accuracy: 0.5218 - loss: 0.6591
1/1 ----- 0s 66ms/step
15/15 ----- 1s 46ms/step - accuracy: 0.5558 - loss: 0.6667 - val_accuracy: 0.5517 - val_loss: 0.6881
Epoch 14/20
4/4 ----- 0s 6ms/step p - accuracy: 0.5251 - loss: 0.6976
1/1 ----- 0s 64ms/step
15/15 ----- 1s 46ms/step - accuracy: 0.5538 - loss: 0.6876 - val_accuracy: 0.5517 - val_loss: 0.6882
Epoch 15/20
4/4 ----- 0s 6ms/step p - accuracy: 0.5137 - loss: 0.6858
1/1 ----- 0s 64ms/step
15/15 ----- 1s 43ms/step - accuracy: 0.5487 - loss: 0.6822 - val_accuracy: 0.5172 - val_loss: 0.6868
Epoch 16/20
4/4 ----- 0s 5ms/step p - accuracy: 0.5399 - loss: 0.6563
1/1 ----- 0s 66ms/step
15/15 ----- 1s 44ms/step - accuracy: 0.5542 - loss: 0.6635 - val_accuracy: 0.5517 - val_loss: 0.6876
Epoch 17/20
4/4 ----- 0s 6ms/step p - accuracy: 0.6219 - loss: 0.668
1/1 ----- 0s 67ms/step
15/15 ----- 1s 44ms/step - accuracy: 0.6164 - loss: 0.6698 - val_accuracy: 0.5862 - val_loss: 0.6838
Epoch 18/20
4/4 ----- 0s 6ms/step p - accuracy: 0.5672 - loss: 0.6456
1/1 ----- 0s 64ms/step
15/15 ----- 1s 43ms/step - accuracy: 0.5487 - loss: 0.6577 - val_accuracy: 0.5172 - val_loss: 0.6852
Epoch 19/20
4/4 ----- 0s 6ms/step ep - accuracy: 0.6452 - loss: 0.654
1/1 ----- 0s 65ms/step
15/15 ----- 1s 49ms/step - accuracy: 0.6284 - loss: 0.6578 - val_accuracy: 0.5172 - val_loss: 0.6868
Epoch 20/20
4/4 ----- 0s 6ms/step p - accuracy: 0.5823 - loss: 0.668
1/1 ----- 0s 59ms/step
15/15 ----- 1s 44ms/step - accuracy: 0.5758 - loss: 0.6715 - val_accuracy: 0.5172 - val_loss: 0.6851

```

More detailed precision graph



```

[286]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import Callback

daily_data = pd.read_csv("XAUUSD_Daily_Ask_2024.01.01_2024.06.30.csv")
daily_data['Time (UTC)'] = pd.to_datetime(daily_data['Time (UTC)'])
daily_data.set_index('Time (UTC)', inplace=True)
daily_data = daily_data[['Close']]

daily_data['Target'] = (daily_data['Close'].shift(-1) > daily_data['Close']).astype(int)
daily_data.dropna(inplace=True)
daily_data['Close'] = (daily_data['Close'] - daily_data['Close'].mean()) / daily_data['Close'].std()

sequence_length = 10
X, y = [], []
for i in range(len(daily_data) - sequence_length):
    X.append(daily_data['Close'].iloc[i:i + sequence_length].values)
    y.append(daily_data['Target'].iloc[i + sequence_length])
X = np.array(X)
y = np.array(y)

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_val = X_val.reshape((X_val.shape[0], X_val.shape[1], 1))

class AccuracyCallback(Callback):
    def __init__(self):
        self.accuracy = []

    def on_epoch_end(self, epoch, logs=None):
        self.accuracy.append(logs['val_accuracy'])

```

```

model = Sequential([
    LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'),
    Dense(1, activation='sigmoid')
])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

accuracy_callback = AccuracyCallback()
model.fit(X_train, y_train, epochs=15, batch_size=8, verbose=1, validation_data=(X_val, y_val), callbacks=[accuracy_callback])


plt.plot(accuracy_callback.accuracy)
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Training accuracy graph')
plt.show()

```

```

Epoch 1/15      8s 86ms/step - accuracy: 0.6213 - loss: 0.6871 - val_accuracy: 0.5172 - val_loss: 0.6905
Epoch 2/15      0s 14ms/step - accuracy: 0.5919 - loss: 0.6855 - val_accuracy: 0.5172 - val_loss: 0.6887
Epoch 3/15      0s 14ms/step - accuracy: 0.5701 - loss: 0.6864 - val_accuracy: 0.5172 - val_loss: 0.6883
Epoch 4/15      0s 14ms/step - accuracy: 0.6151 - loss: 0.6628 - val_accuracy: 0.5172 - val_loss: 0.6899
Epoch 5/15      0s 14ms/step - accuracy: 0.5518 - loss: 0.6906 - val_accuracy: 0.5172 - val_loss: 0.6900
Epoch 6/15      0s 14ms/step - accuracy: 0.5695 - loss: 0.6979 - val_accuracy: 0.5172 - val_loss: 0.6900
Epoch 7/15      0s 16ms/step - accuracy: 0.5799 - loss: 0.6797 - val_accuracy: 0.5517 - val_loss: 0.6918
Epoch 8/15      0s 16ms/step - accuracy: 0.5804 - loss: 0.7042 - val_accuracy: 0.5517 - val_loss: 0.6901
Epoch 9/15      0s 17ms/step - accuracy: 0.5763 - loss: 0.6813 - val_accuracy: 0.5172 - val_loss: 0.6892
Epoch 10/15     0s 13ms/step - accuracy: 0.5349 - loss: 0.7019 - val_accuracy: 0.5517 - val_loss: 0.6892
Epoch 11/15     0s 14ms/step - accuracy: 0.5925 - loss: 0.6821 - val_accuracy: 0.5517 - val_loss: 0.6887
Epoch 12/15     0s 13ms/step - accuracy: 0.5529 - loss: 0.6778 - val_accuracy: 0.5172 - val_loss: 0.6890
Epoch 13/15     0s 16ms/step - accuracy: 0.5866 - loss: 0.6645 - val_accuracy: 0.5172 - val_loss: 0.6893
Epoch 14/15     0s 17ms/step - accuracy: 0.6401 - loss: 0.6639 - val_accuracy: 0.5517 - val_loss: 0.6882
Epoch 15/15     0s 13ms/step - accuracy: 0.5553 - loss: 0.6868 - val_accuracy: 0.5517 - val_loss: 0.6869

```



```
[288]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import Callback

daily_data = pd.read_csv("XAUUSD_Daily_Ask_2024.01.01_2024.06.30.csv")
daily_data["Time (UTC)"] = pd.to_datetime(daily_data["Time (UTC)"])
daily_data.set_index("Time (UTC)", inplace=True)
daily_data = daily_data[["Close"]]

daily_data["Target"] = (daily_data["Close"].shift(-1) > daily_data["Close"]).astype(int)
daily_data.dropna(inplace=True)
daily_data["Close"] = (daily_data["Close"] - daily_data["Close"].mean()) / daily_data["Close"].std()

sequence_length = 10
X, y = [], []
for i in range(len(daily_data) - sequence_length):
    X.append(daily_data["Close"].iloc[i:i + sequence_length].values)
    y.append(daily_data["Target"].iloc[i + sequence_length])
X = np.array(X)
y = np.array(y)

X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2, random_state=42)
X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))
X_val = X_val.reshape((X_val.shape[0], X_val.shape[1], 1))

class AccuracyCallback(Callback):
    def __init__(self):
        self.train_accuracy = []
        self.val_accuracy = []

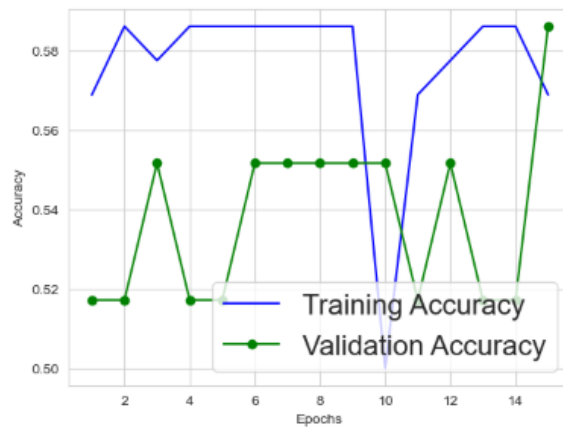
    def on_epoch_end(self, epoch, logs=None):
        self.train_accuracy.append(logs['accuracy'])
        self.val_accuracy.append(logs['val_accuracy'])

model = Sequential([
    LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'),
    Dense(1, activation='sigmoid') # Binary classification
])
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

accuracy_callback = AccuracyCallback()
model.fit(X_train, y_train, epochs=15, batch_size=8, verbose=1, validation_data=(X_val, y_val), callbacks=[accuracy_callback])

epochs = range(1, len(accuracy_callback.train_accuracy) + 1)
plt.plot(epochs, accuracy_callback.train_accuracy, label='Training Accuracy', color='blue')
plt.plot(epochs, accuracy_callback.val_accuracy, label='Validation Accuracy', color='green', marker='o')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

```
Epoch 1/15
15/15 — 9s 99ms/step - accuracy: 0.6340 - loss: 0.6865 - val_accuracy: 0.5172 - val_loss: 0.6900
Epoch 2/15
15/15 — 0s 14ms/step - accuracy: 0.5411 - loss: 0.6914 - val_accuracy: 0.5172 - val_loss: 0.6886
Epoch 3/15
15/15 — 0s 14ms/step - accuracy: 0.5361 - loss: 0.6769 - val_accuracy: 0.5517 - val_loss: 0.6877
Epoch 4/15
15/15 — 0s 14ms/step - accuracy: 0.5294 - loss: 0.6881 - val_accuracy: 0.5172 - val_loss: 0.6884
Epoch 5/15
15/15 — 0s 14ms/step - accuracy: 0.5850 - loss: 0.6820 - val_accuracy: 0.5172 - val_loss: 0.6890
Epoch 6/15
15/15 — 0s 15ms/step - accuracy: 0.6284 - loss: 0.6660 - val_accuracy: 0.5517 - val_loss: 0.6928
Epoch 7/15
15/15 — 0s 16ms/step - accuracy: 0.6018 - loss: 0.6597 - val_accuracy: 0.5517 - val_loss: 0.6906
Epoch 8/15
15/15 — 0s 15ms/step - accuracy: 0.6017 - loss: 0.6935 - val_accuracy: 0.5517 - val_loss: 0.6886
Epoch 9/15
15/15 — 0s 13ms/step - accuracy: 0.6411 - loss: 0.6592 - val_accuracy: 0.5517 - val_loss: 0.6882
Epoch 10/15
15/15 — 0s 14ms/step - accuracy: 0.5875 - loss: 0.6570 - val_accuracy: 0.5517 - val_loss: 0.6870
Epoch 11/15
15/15 — 0s 13ms/step - accuracy: 0.5645 - loss: 0.6803 - val_accuracy: 0.5172 - val_loss: 0.6855
Epoch 12/15
15/15 — 0s 13ms/step - accuracy: 0.5788 - loss: 0.6932 - val_accuracy: 0.5517 - val_loss: 0.6872
Epoch 13/15
15/15 — 0s 14ms/step - accuracy: 0.5461 - loss: 0.6774 - val_accuracy: 0.5172 - val_loss: 0.6867
Epoch 14/15
15/15 — 0s 14ms/step - accuracy: 0.5766 - loss: 0.6606 - val_accuracy: 0.5172 - val_loss: 0.6861
Epoch 15/15
15/15 — 0s 13ms/step - accuracy: 0.5700 - loss: 0.6740 - val_accuracy: 0.5862 - val_loss: 0.6840
```



[336]: # Calculate the prediction vector

```
# Verify and reshape X_test
print("Original X_test shape:", X_test.shape)
if len(X_test.shape) == 2:
    X_test = X_test.reshape((X_test.shape[0], X_test.shape[1], 1))
print("Reshaped X_test shape:", X_test.shape)

# Check model input shape
print("Expected input shape:", model.input_shape)

# Make predictions
pred = model.predict(X_test)
print("Predictions:", pred)
```

```
Original X_test shape: (29, 10, 1)
Reshaped X_test shape: (29, 10, 1)
Expected input shape: (None, 10, 1)
1/1 — 0s 116ms/step
Predictions: [[0.5556013 ]
 [0.4996255 ]
 [0.6276852 ]
 [0.63105965]
 [0.6306144 ]
 [0.49330345]
 [0.5002903 ]
 [0.4950856 ]
 [0.49369597]
 [0.63282484]
 [0.63177073]
 [0.5741288 ]
 [0.5018753 ]
 [0.50046986]
 [0.5219916 ]
 [0.6392977 ]
 [0.49119493]
 [0.50059724]
 [0.50257075]
 [0.50231224]
 [0.500361 ]
 [0.50128865]
 [0.63047624]
 [0.5686562 ]
 [0.63633466]
 [0.6273813 ]
 [0.5925585 ]
 [0.5055052 ]
 [0.49101394]]
```

```
[296]: print(pred)
```

```
[[0.5556013 ]
 [0.4996255 ]
 [0.6276852 ]
 [0.63105965]
 [0.6306144 ]
 [0.49330345]
 [0.5002903 ]
 [0.4950856 ]
 [0.49369597]
 [0.63282484]
 [0.63177073]
 [0.5741288 ]
 [0.5018753 ]
 [0.50046986]
 [0.5219916 ]
 [0.6392977 ]
 [0.49119493]
 [0.50059724]
 [0.50257075]
 [0.50231224]
 [0.500361 ]
 [0.50128805]
 [0.63047624]
 [0.5686562 ]
 [0.63633466]
 [0.6273813 ]
 [0.5925585 ]
 [0.5055052 ]
 [0.49101394]]
```

```
[298]: len(pred)
```

```
[298]: 29
```

```
[ ]:
```

```
[300]: import random
```

```
pred = model.predict(X_test)
```

```
# Check: we take a random element random.randint() and look: what is the difference between test and predict
```

```
n_rec = random.randint(0, X_test.shape[0])
print(n_rec)
```

```
print("Predicted probability:", pred[n_rec], ", right answer:", y_test[n_rec])
```

```
1/1 ————— 0s 83ms/step
```

```
5
```

```
Predicted probability: [0.49330345] , right answer: 1
```

```
[326]: classes=['0 is Flat', '1 is Trend']

index = random.randint(0, y_test.shape[0])
print('Right answer: ', y_test[index])

x = X_test[index]
x = np.expand_dims(x, axis=0)

prediction = model.predict(x)
sample = x

ans = round(float(prediction))


fig = plt.figure(figsize=(5,3))

ax = fig.add_subplot(1, 2, 2)
bar_list = ax.bar(np.arange(1), prediction[0], align='center')
if ans == y_test[index]:
    bar_list[0].set_color('g')
else:
    bar_list[0].set_color('r')

ax.set_xticks(np.arange(1))
ax.set_xlim([-1, 1])
ax.grid('on')

plt.show()
print('Predicted answer: {}'.format(classes[ans]), "\n ")
print('Right answer: {}'.format(classes[y_test[index].astype(int)]) )
print(classes)

Right answer: 1
1/1 ————— 0s 80ms/step
```



Predicted answer: 1 is Trend

Right answer: 1 is Trend
 ['0 is Flat', '1 is Trend']

Lab 11

Lab Logbook Requirement:

1. Create and train your own LSTM model
2. Add all the LSTM's Error metrics: Accuracy, Precision, Recall, F1-Score and AUC to the final histogram "ML Models performance..."

```
[152]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score
import matplotlib.pyplot as plt

credit_data = pd.read_csv("credit_risk_dataset.csv")

credit_data['person_home_ownership'] = LabelEncoder().fit_transform(credit_data['person_home_ownership'])
credit_data['loan_intent'] = LabelEncoder().fit_transform(credit_data['loan_intent'])
credit_data['cb_person_default_on_file'] = LabelEncoder().fit_transform(credit_data['cb_person_default_on_file'])

credit_data.fillna(credit_data.median(), inplace=True)

X = credit_data.drop('loan_status', axis=1).values
y = credit_data['loan_status'].values

scaler = StandardScaler()
X = scaler.fit_transform(X)

X = X.reshape((X.shape[0], 1, X.shape[1]))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = Sequential([
    LSTM(50, input_shape=(X_train.shape[1], X_train.shape[2]), activation='relu'),
    Dense(1, activation='sigmoid')
])

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=10, batch_size=32, verbose=1, validation_data=(X_test, y_test))

y_pred = (model.predict(X_test) > 0.5).astype(int)

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, model.predict(X_test))

metrics = {
    'Accuracy': accuracy,
    'Precision': precision,
    'Recall': recall,
    'F1-Score': f1,
    'AUC': auc
}

plt.bar(metrics.keys(), metrics.values())
plt.title("ML Models Performance")
plt.ylabel("Score")
plt.show()
```

Epoch 1/10
815/815 18s 7ms/step - accuracy: 0.7968 - loss: 0.4926 - val_accuracy: 0.8373 - val_loss: 0.3728
Epoch 2/10
815/815 5s 5ms/step - accuracy: 0.8531 - loss: 0.3549 - val_accuracy: 0.8558 - val_loss: 0.3509
Epoch 3/10
815/815 4s 5ms/step - accuracy: 0.8613 - loss: 0.3392 - val_accuracy: 0.8599 - val_loss: 0.3430
Epoch 4/10
815/815 4s 5ms/step - accuracy: 0.8677 - loss: 0.3317 - val_accuracy: 0.8601 - val_loss: 0.3385
Epoch 5/10
815/815 4s 5ms/step - accuracy: 0.8646 - loss: 0.3358 - val_accuracy: 0.8665 - val_loss: 0.3320
Epoch 6/10
815/815 5s 5ms/step - accuracy: 0.8702 - loss: 0.3224 - val_accuracy: 0.8690 - val_loss: 0.3287
Epoch 7/10
815/815 5s 6ms/step - accuracy: 0.8682 - loss: 0.3284 - val_accuracy: 0.8668 - val_loss: 0.3277
Epoch 8/10
815/815 4s 5ms/step - accuracy: 0.8719 - loss: 0.3188 - val_accuracy: 0.8700 - val_loss: 0.3238
Epoch 9/10
815/815 5s 6ms/step - accuracy: 0.8742 - loss: 0.3157 - val_accuracy: 0.8682 - val_loss: 0.3220
Epoch 10/10
815/815 5s 5ms/step - accuracy: 0.8783 - loss: 0.3084 - val_accuracy: 0.8752 - val_loss: 0.3207
204/204 2s 7ms/step
204/204 1s 4ms/step

